

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import numpy as np
from scipy.stats import iqr
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

df = pd.read_csv("marketing_campaign.csv", sep="\t")
df.head()
```

Out[1]:

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer	Recency	MntWines	...	NumWebVisit
0	5524	1957	Graduation	Single	58138.0	0	0	04-09-2012	58	635	...	
1	2174	1954	Graduation	Single	46344.0	1	1	08-03-2014	38	11	...	
2	4141	1965	Graduation	Together	71613.0	0	0	21-08-2013	26	426	...	
3	6182	1984	Graduation	Together	26646.0	1	0	10-02-2014	26	11	...	
4	5324	1981	PhD	Married	58293.0	1	0	19-01-2014	94	173	...	

5 rows × 29 columns

```
In [2]: df.dtypes
```

Out[2]:

ID	int64
Year_Birth	int64
Education	object
Marital_Status	object
Income	float64
Kidhome	int64
Teenhome	int64
Dt_Customer	object
Recency	int64
MntWines	int64
MntFruits	int64
MntMeatProducts	int64
MntFishProducts	int64
MntSweetProducts	int64
MntGoldProds	int64
NumDealsPurchases	int64
NumWebPurchases	int64
NumCatalogPurchases	int64
NumStorePurchases	int64
NumWebVisitsMonth	int64
AcceptedCmp3	int64
AcceptedCmp4	int64
AcceptedCmp5	int64
AcceptedCmp1	int64
AcceptedCmp2	int64
Complain	int64
Z_CostContact	int64
Z_Revenue	int64
Response	int64
dtype:	object

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                     2240 non-null   int64
1   Year_Birth                           2240 non-null   int64
2   Education                             2240 non-null   object
3   Marital_Status                       2240 non-null   object
4   Income                               2216 non-null   float64
5   Kidhome                              2240 non-null   int64
6   Teenhome                             2240 non-null   int64
7   Dt_Customer                          2240 non-null   object
8   Recency                              2240 non-null   int64
9   MntWines                             2240 non-null   int64
10  MntFruits                             2240 non-null   int64
11  MntMeatProducts                       2240 non-null   int64
12  MntFishProducts                       2240 non-null   int64
13  MntSweetProducts                     2240 non-null   int64
14  MntGoldProds                         2240 non-null   int64
15  NumDealsPurchases                    2240 non-null   int64
16  NumWebPurchases                      2240 non-null   int64
17  NumCatalogPurchases                 2240 non-null   int64
18  NumStorePurchases                   2240 non-null   int64
19  NumWebVisitsMonth                   2240 non-null   int64
20  AcceptedCmp3                        2240 non-null   int64
21  AcceptedCmp4                        2240 non-null   int64
22  AcceptedCmp5                        2240 non-null   int64
23  AcceptedCmp1                        2240 non-null   int64
24  AcceptedCmp2                        2240 non-null   int64
25  Complain                             2240 non-null   int64
26  Z_CostContact                       2240 non-null   int64
27  Z_Revenue                           2240 non-null   int64
28  Response                             2240 non-null   int64
dtypes: float64(1), int64(25), object(3)
memory usage: 507.6+ KB
```

```
In [4]: df["TotalAmountSpent"] = df["MntFishProducts"] + df["MntFruits"] + df["MntGoldProds"] + df["MntSweetProducts"] +
```

Univariate analysis Univariate analysis entails evaluating a single feature in order to get insights about it. So, the initial step in performing EDA is to undertake univariate analysis, which includes evaluating descriptive or summary statistics about the feature.

For example you might check a feature distribution, proportion of a feature, and so on.

In our case, we will check the distribution of customer's ages in the dataset. We can do that by typing the following:

```
In [6]: print(df.columns.tolist())
```


```
['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income', 'Kidhome', 'Teenhome', 'Dt_Customer', 'Recency', 'MntWines', 'MntFruits', 'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts', 'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases', 'NumWebVisitsMonth', 'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1', 'AcceptedCmp2', 'Complain', 'Z_CostContact', 'Z_Revenue', 'Response', 'TotalAmountSpent']
```

```
In [11]: df.columns = df.columns.str.strip()
```

```
In [13]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from datetime import datetime

# Assuming 'Year_Birth' is the year of birth
current_year = datetime.now().year
df['Age'] = current_year - df['Year_Birth']

# Plotting the histogram
sns.histplot(data=df, x="Age", bins=list(range(10, 150, 10)))
plt.title("Distribution of Customer's Age")
plt.show()
```



Bivariate Analysis After you've performed univariate analysis on all your feature of interest, the next step is to perform bivariate analysis. This involves comparing two attributes at the same time.

Bivariate analysis entails determining the correlation between two features, for example.

In our case, some of the bivariate analysis we'll perform in the project include observing the average total spent across different client age groups, determining a correlation between customer income and total amount spent, and so on, as shown below.

For example, in our case we want to check the relationship between a Customer's Income and TotalAmountSpent. We can do that by typing the following:

```
In [15]: fig = px.scatter(data_frame=df, x="Income",
                        y="TotalAmountSpent",
                        title="Relationship Between Customer's Income and Total Amount Spent",
                        height=500,
                        color_discrete_sequence = px.colors.qualitative.G10[1:])
fig.show()
```

We can see from the above analysis that as the Income increases so does the TotalAmountSpent. So from the analysis we can postulate that Income is one of key factor that determines how much a customer might spend.

Multivariate Analysis After you've completed univariate (analysis of single feature) and bivariate (analysis of two features) analysis, the last phase of EDA is to perform Multivariate Analysis.

Multivariate Analysis consists of understanding the relationship between two or more variables.

In our project, one of the multivariate analysis we'll do is to understand the relationship between Income, TotalAmountSpent, and Customer's Education.

```
In [17]: fig = px.scatter(
          data_frame=df,
          x = "Income",
          y= "TotalAmountSpent",
          title = "Relationship between Income VS Total Amount Spent Based on Education",
          color = "Education",
          height=500
        )
fig.show()
```

How to Build the Segmentation Model After we've finished our analysis, the next step is to create the model that will segment the customers. KMeans is the model we'll use. It is a popular segmentation model that is also quite effective.

The KMeans model is an unsupervised machine learning model that works by simply splitting N observations into K numbers of clusters. The observations are grouped into these clusters based on how close they are to the mean of that cluster, which is commonly referred to as centroids.

When you fit the features into the model and specify the number of clusters or segments you want, KMeans will output the cluster label to which each observation in the feature belongs.

Let's talk about the features you might want to fit into a KMeans model. There are no limits to the number of features you can use to build a Customer segmentation model – but in my opinion, fewer's better. This is because you will be able to grasp and interpret the outcomes of each segment more easily and clearly with fewer features.

In our scenario, we will first construct the KMeans model with two features and then build the final model with three features. But, before we get started, let's go over the KMeans assumptions, which are as follows:

The features must be numerical.

The features you're fitting into KMeans must be normally distributed. This is because KMeans (since it calculates average distance) is affected by outliers (values that deviate a lot from the others). As a result, any skewed feature must be changed in order to be normally distributed. Fortunately, we can use Numpy's logarithm transformation package `np.log()`

The features must also be of the same scale. For this, we'll use the Scikit-learn `StandardScaler()` module.

We'll design our KMeans model now that we've grasped the main concept. So, for our first model, we'll use the Income and TotalAmountSpent features.

To begin, because the Income feature has missing values, we will fill it with the median number.

```
In [19]: df['Age'] = current_year - df['Year_Birth']
```

```

In [20]: df = df.copy() # Ensure you're working with a copy of the DataFrame
df['Age'] = current_year - df['Year_Birth']

In [23]: df['Income'] = df['Income'].fillna(df['Income'].median())

In [25]: #After that, we'll assign the features we want to work with, Income and TotalAmountSpent, to a variable called data

In [26]: data = df[["Income", "TotalAmountSpent"]]

In [27]: ##Once that's done we will transform features and save the result into a variable called data_log.

In [28]: df_log = np.log(data)

In [30]: ##Then we will scale the result using Scikit-learn StandardScaler():

std_scaler = StandardScaler()
df_scaled = std_scaler.fit_transform(df_log)

```

Once that's done we can then build the model. So the KMeans model requires two parameters. The first is `random_state` and the second one is `n_clusters` where:

`n_clusters` represents the number of clusters or segments to be derived from KMeans.

`random_state`: is required for reproducible results.

```


In [34]: errors = []
for k in range(1, 11):
    model = KMeans(n_clusters=k, random_state=42)
    model.fit(df_scaled)
    errors.append(model.inertia_)

plt.title('The Elbow Method')
plt.xlabel('k')
plt.ylabel('Error of Cluster')

# Corrected pointplot syntax
sns.pointplot(x=list(range(1, 11)), y=errors)

plt.show()

```



Let's summarize what the above code does. We specified the number of clusters to experiment with, which is in the `range(1, 11)`. Then we fit the features on those clusters and added the error to the list we created before above.

Following that, we plot the error for each cluster. The diagram shows that the cluster that creates the elbow is three. So three clusters is the best value for our model. As a result, we will build the KMeans model utilizing three clusters.

```

In [35]: model = KMeans(n_clusters = 3, random_state=42)
model.fit(df_scaled)

```

```
Out[35]: KMeans
KMeans(n_clusters=3, random_state=42)
```

```
In [37]: ##Now we've built our model. The next thing will be to assign the cluster label for each observation. So we will
data = data.assign(ClusterLabel = model.labels_)
```

```
In [38]: data.groupby("ClusterLabel")[["Income", "TotalAmountSpent"]].median()
```

```
Out[38]:
```

	Income	TotalAmountSpent
ClusterLabel		
0	32765.0	57.0
1	157243.0	107.0
2	65203.0	934.0

We can see that there is a trend within the clusters:

Cluster 0 translates to customers who earn less and spend less.

Cluster 1 represent customers that earn more and spend more.

Cluster 2 represents customers that earn moderate and spend moderate.

```
In [39]: fig = px.scatter(
    data_frame=data,
    x = "Income",
    y= "TotalAmountSpent",
    title = "Relationship between Income VS Total Amount Spent",
    color = "ClusterLabel",
    height=500
)
fig.show()
```

```
In [40]: data = df[["Age", "Income", "TotalAmountSpent"]]
df_log = np.log(data)
std_scaler = StandardScaler()
df_scaled = std_scaler.fit_transform(df_log)
```

```
In [43]: # Assuming 'data' contains the 'ClusterLabel' column
data = data.assign(ClusterLabel=model.labels_)

# Group by the 'ClusterLabel' and calculate mean and median for the specified columns
result = data.groupby("ClusterLabel").agg({"Age": "mean", "Income": "median", "TotalAmountSpent": "median"}).reset_index()
```

```
In [45]: model = KMeans(n_clusters=3, random_state=42)
```

```

model.fit(df_scaled)

data = data.assign(ClusterLabel= model.labels_)


result = data.groupby("ClusterLabel").agg({"Age": "mean", "Income": "median", "TotalAmountSpent": "median"}).round

```

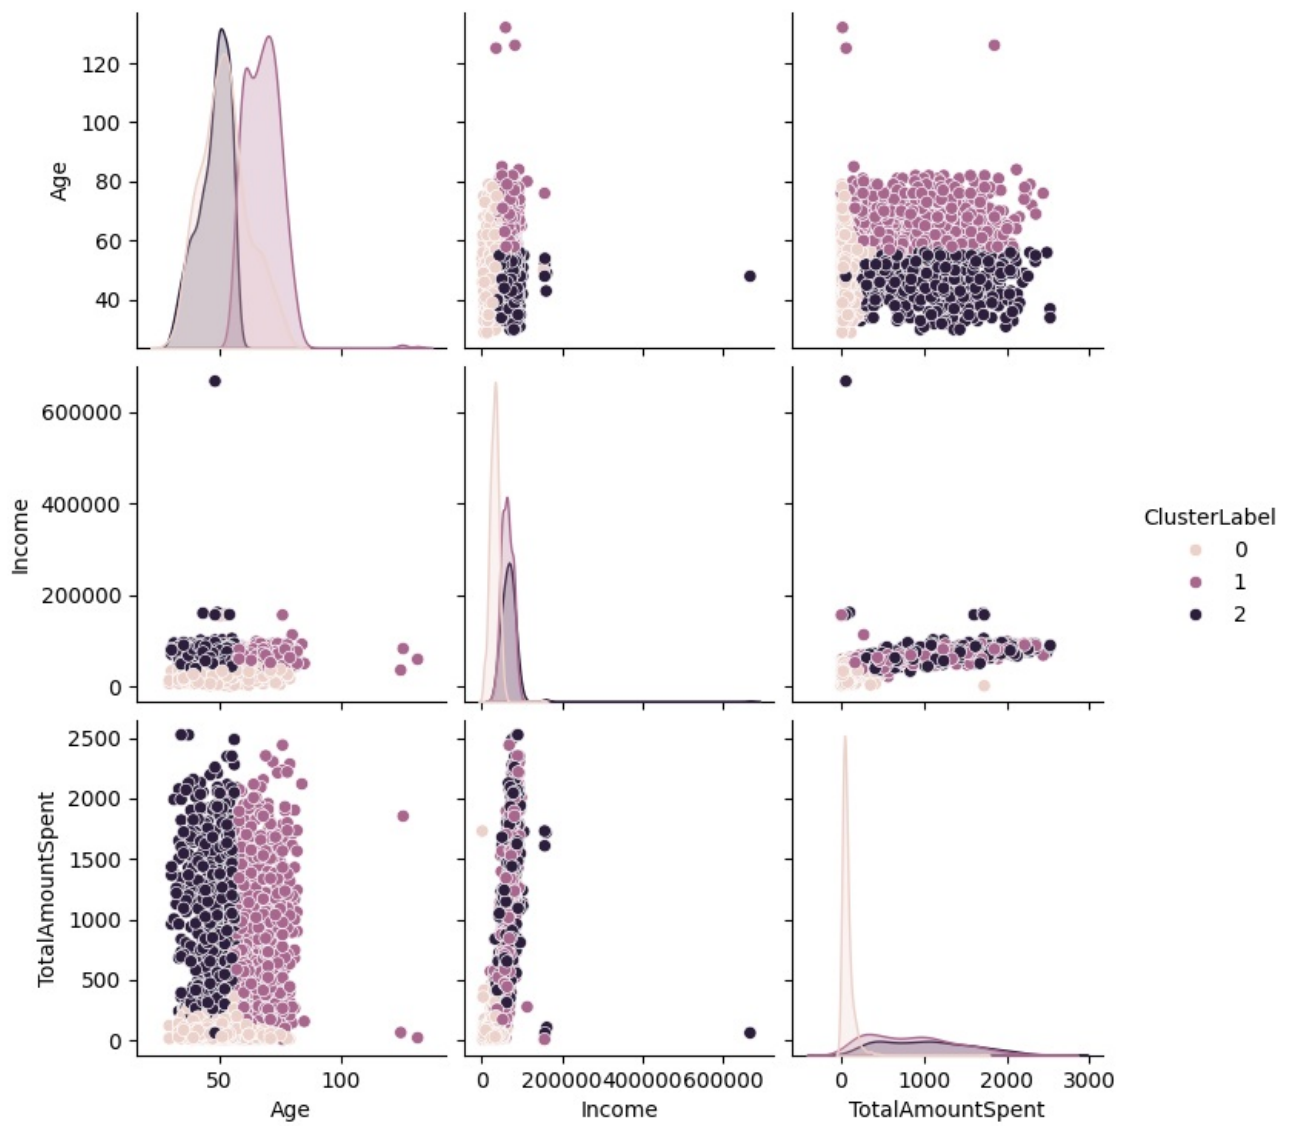
In [48]: `print(result)`

ClusterLabel	Age	Income	TotalAmountSpent
0	52.0	31801.0	54.0
1	68.0	62820.0	825.0
2	47.0	67384.0	1001.0

In [49]: `cluster_centers = model.cluster_centers_`
`plt.scatter(df_scaled[:, 0], df_scaled[:, 1], c=model.labels_, cmap='viridis')`
`plt.scatter(cluster_centers[:, 0], cluster_centers[:, 1], s=200, c='red', marker='X') # Red X for centroids`
`plt.title('KMeans Clusters and Centroids')`
`plt.show()`



In [50]: `import seaborn as sns`
`data['ClusterLabel'] = model.labels_`
`sns.pairplot(data, hue='ClusterLabel', vars=['Age', 'Income', 'TotalAmountSpent'])`
`plt.show()`



Cluster 0 depicts young customers that earn a lot and also spend a lot.

Cluster 1 translates to older customers that earn a lot and also spend a lot.

Cluster 2 depicts young customers that earn less and also spend less.

```
In [54]: fig = px.scatter_3d(data_frame=data, x="Income",
                             y="TotalAmountSpent", z="Age", color="ClusterLabel", height=550,
                             title = "Visualizing Cluster Result Using 3 Features")
fig.show()
```


Conclusion In this tutorial, you learnt how to build a customer segmentation model. There are a lot of features we didn't touch on in this article. But I suggest that you experiment with it and create customer segmentation models using different features.

In []: